

Alpaca Options Trading Agents

Miramar Labs · Alpaca AI Trading Agents Hackathon · lablab.ai · Sep 2026 · github.com/miramar-labs/alpaca-options-trading-agents

Paper account **PA3MAA2HJLUG** · \$100k

A three-agent trading floor that autonomously trades US-equity **options**. An **Analyst** sets the daily universe; a **Dealer** produces a BUY / HOLD / SELL signal per symbol every 600 seconds and hands each non-HOLD to an **Alpaca MCP tool-calling agent** that browses the live option chain and picks one specific contract; a **Floor Broker** re-validates that pick server-side and executes it. It runs live and continuously on a Kubernetes cluster against a dedicated Alpaca paper account — not a backtest or a one-off demo.

Analyst (CronJob 08:55 ET) → portfolio ConfigMap → Dealer (poll loop, 600s) → BUY/HOLD/SELL (LLM) → non-HOLD → Alpaca MCP loop (chain → contract) → Floor Broker (FastAPI) → Alpaca: re-validate, place order, synthetic SL/TP/DTE exit

AI LOGIC

- **Analyst** — a LangGraph pipeline: screens Alpaca movers plus a fixed large-cap list, pulls news and TAAPI technical indicators, reads back its own recent track record from Postgres, and asks an LLM for the day's watchlist with a budget and rationale per symbol.
- **Dealer signal** — per symbol, technical indicators and OHLCV-derived market structure go to an LLM for a BUY / HOLD / SELL call plus a confidence score.
- **Contract selection** (the centrepiece) — `mcp_options.py` launches `alpaca-mcp-server` over stdio (`langchain-mcp-adapters`) with a **read-only** toolset. Given the underlying, the direction (call = bullish, put = bearish; always long) and the config DTE / delta windows, the LLM drives a bounded **≤6-round** tool-calling loop over the live chain, quotes and Greeks and returns a structured pick: strike, expiration, delta, premium, reasoning.
- **Validate** → **reconcile** → **re-validate** — reject a wrong-type or never-seen-in-chain contract; overwrite strike / expiration / delta / premium with what the chain actually showed (only the reasoning text is trusted); re-check server-side in Floor Broker against live windows and a fresh ask quote before the order goes in.
- **Local inference only** — Ollama serving `qwen2.5:32b-instruct-q4_K_M` for every decision; no external LLM API. If the structured-output step ever returns empty, a deterministic `_fallback_pick` selects from the same fetched chain rows under the same gates — a miss degrades to “code skims the chain like a human would,” never to no trade.

RISK GATES

Eight layers sit between “the LLM said BUY” and an order reaching Alpaca — each one deterministic, none of them an LLM call:

- confidence threshold on the signal
- macro-event blackout calendar (FOMC / CPI / jobs / PCE)
- same-symbol stop-loss cooldown
- portfolio win-rate throttle
- daily profit-target / loss-limit halt
- operator kill switch — flipped in config, no redeploy
- DTE / delta window re-validation
- hard notional cap, enforced against the live market rather than the model's claimed price

Once a position is open, Alpaca has no server-side bracket order for options, so Floor Broker runs its own 30-second poll loop that force-closes on a **50% loss**, a **175% gain**, or expiration dropping to **3 days** out — whichever comes first, and regardless of whether opening new positions is currently gated.

ALPACA INFRASTRUCTURE IMPLEMENTATION

- `alpaca-py` for trading, market data, news and the options chain / quotes / Greeks — all against the dedicated paper account (`PA3MAA2HJLUG`).
- `alpaca-mcp-server`, the official MCP server, run locally as a stdio subprocess per Dealer poll with a read-only toolset — the agent can inspect the chain but can never place an order.
- **Floor Broker** is the only component with write access to the Trading API: a stateless FastAPI service that submits the option order over plain REST, then runs three background poll loops — fill detection, synthetic SL/TP/DTE exit, and restart recovery (rebuilding all in-memory tracking from Alpaca and Postgres if the pod restarts mid-position).
- **One NVIDIA DGX Spark** hosts everything: k3s (four workloads in namespace `alpaca-options-trader`), Ollama, a shared Postgres instance (append-only audit trail, track record, DB-backed reconciliation), and a self-hosted GitHub Actions runner doing the whole CI/CD path (push → test + lint → build four images → rolling deploy).
- `config.yaml` is fetched fresh from GitHub every 60 seconds by every pod, so risk thresholds, blackout dates, even the LLM model are a plain `git push` — no rebuild, no redeploy.

RESULTS — DAY-2 SNAPSHOT (2 SEP 2026)

Activity. 114 Dealer cycles (97 HOLD / 15 BUY / 2 SELL). The 15 BUY signals produced **7 option positions**, all filled within half a minute; the rest were stopped by the risk layer — one sub-0.6 confidence, several duplicates of an open position or an order in flight, three with no chain contract passing the DTE / delta / OI / volume gates.

Open book. 5 long calls, 2 long puts, ~\$10.0k premium deployed, entry deltas 0.43–0.53, 15–45 DTE. Account equity **\$100,824** (+0.8% vs the \$100k open); aggregate unrealized P/L **+\$825**, 5 of 7 positions green. Best an FRVO put at +32%; worst an MSFT call at -37%. No position has closed — no realized P/L or win rate yet.

Gates observed firing. The daily profit-target halt tripped three times once intraday P/L cleared +\$1,000; the confidence, duplicate and contract-quality gates each rejected a real signal; no synthetic stop-loss / take-profit / DTE close has triggered. All 7 day-1–2 contracts were chosen by the deterministic `_fallback_pick` path — the MCP loop fetched live chain data every cycle, only the closing structured-output step came back empty on `qwen3.6:35b-a3b`. Root-caused to that model's low active-parameter count and fixed inside the window by switching to `qwen2.5:32b` (real-path tested; see `docs/models.md`); selection now runs on the model. Live P/L and model badges are in the README.